

PyTorch for .NET Developers with Azure Machine Learning

Adnan Masood, PhD

Microsoft Azure
+ AI Conference

CO-PRODUCED BY
Microsoft & DEVintersection





About the Speaker



Adnan Masood, Ph.D. is an Artificial Intelligence and Machine Learning researcher, software architect, and Microsoft MVP (Most Valuable Professional) for Artificial Intelligence. As Chief Architect of AI and Machine Learning, at UST Global, he collaborates with Stanford Artificial Intelligence Lab, and MIT AI Lab for building enterprise solutions

Author of Amazon bestseller in programming languages, "**Functional Programming with F#**", Dr. Masood teaches Data Science at Park University, and has taught Windows Communication Foundation (WCF) courses at the University of California, San Diego. He is a regular speaker to various academic and technology conferences (WICT, DevIntersection, IEEE-HST, IASA, and DevConnections), local code camps, and user groups. He also volunteers as STEM (Science Technology, Engineering and Math) robotics coach for elementary and middle school students

A strong believer in giving back to the community, Dr. Masood is a co-founder and president of the Pasadena .NET Developers group, co-organizer of Tampa Bay Data Science Group, and Irvine Programmer meetup. His recent talk at Women in Technology Conference (WICT) Denver highlighted the importance of diversity in STEM and technology areas, and was featured by variety of news outlets.





Andrej Karpathy ✓

@karpathy

Follow



I've been using PyTorch a few months now and I've never felt better. I have more energy. My skin is clearer. My eye sight has improved.

11:56 AM - 26 May 2017

399 Retweets **1,532** Likes



PyTorch for .NET Developers with Azure Machine Learning

- PyTorch is a Python-based scientific computing package that uses the power of modern GPUS. It is quickly becoming a preferred deep learning research platform built to provide maximum flexibility and speed using dynamic graph.
- In this talk we will show how Pytorch can be used on Azure using Azure Machine Learning service, Data Science Virtual Machine, Azure Notebooks, and Visual Studio Code Tools for AI. We will train PyTorch models with Azure Machine Learning service, and how we can use the ONNX runtime inside of our .NET applications. ONNX is a open format to represent deep learning models that is supported by various frameworks and tools. This format makes it easier to interoperate between frameworks and to maximize the reach of your hardware optimization investments.

Because writing ML code from scratch sucks (part 1)

```
5 d = size(trainData, 2);
6 step = 0.1;
7 %%
8 m_train = size(trainData, 1);
9
10 %%
11 d1_arr = [1, 5, 10, 15, 25, 50];
12 err_cv_arr = [0, 0, 0, 0, 0, 0];
13 err_train_arr = [0, 0, 0, 0, 0, 0];
14 err_test_arr = [0, 0, 0, 0, 0, 0];
15 %%
16 for i = 1:6
17     d1 = d1_arr(i, 1);
18     w1_init = load(strcat(path_init, "w1_", num2str(d1), ".mat"));
19     w1_init = w1_init.w_1;
20
21     b1_init = load(strcat(path_init, "b1_", num2str(d1), ".mat"));
22     b1_init = b1_init.b1;
23
24     w2_init = load(strcat(path_init, "w2_", num2str(d1), ".mat"));
25     w2_init = w2_init.w_2;
26
27     b2_init = load(strcat(path_init, "b2_", num2str(d1), ".mat"));
28     b2_init = b2_init.b2;
29
30     err_cv_acc = 0;
31     for k = 1:5
32         cv_train = load(strcat(path_cv, num2str(k), "/cv-train.txt"));
33         cv_test = load(strcat(path_cv, num2str(k), "/cv-test.txt"));
34
35         x_cv_train = cv_train;
36         x_cv_train(:, d) = [];
37         y_cv_train = cv_train(:, d);
38         y_cv_train = (y_cv_train + 1) ./ 2;
39         m_cv_train = size(cv_train, 1);
40
41         x_cv_test = cv_test;
42         x_cv_test(:, d) = [];
43         y_cv_test = cv_test(:, d);
44         y_cv_test = (y_cv_test + 1) ./ 2;
45
46         w1 = w1_init;
47         b1 = b1_init;
48         w2 = w2_init;
49         b2 = b2_init;
50
51         for iter = 1:5000
52             z = x_cv_train * w1 + b1;
53             a = 1 ./ (1 + exp(-z));
54             eta = 1 ./ (1 + exp(-(a*w2+b2)));
55             eps = eta - y_cv_train;
56             dz = a * (1 - a) * ...
```

```
58         dw1 = (w2 .* ((eps .* dz) .* x_cv_train)) ./ m_cv_train;
59         db1 = (w2 .* (dz .* eps)) ./ m_cv_train;
60         dw2 = (eps .* a) ./ m_cv_train;
61         db2 = sum(eps) / m_cv_train;
62
63         w1 = w1 - step .* dw1;
64         b1 = b1 - step .* db1;
65         w2 = w2 - step .* dw2;
66         b2 = b2 - step .* db2;
67     end
68
69     pred_a_cv = 1 ./ (1 + exp(-(x_cv_test * w1 + b1)));
70     pred_eta_cv = 1 ./ (1 + exp(-(pred_a_cv * w2 + b2)));
71     pred_y_cv = sign(pred_eta_cv .* 2 - 1);
72     pred_y_cv = (pred_y_cv + 1) ./ 2;
73     err_cv = classification_error(pred_y_cv, y_cv_test);
74
75     err_cv_acc = err_cv_acc + err_cv;
76 end
77 err_cv_arr(i, 1) = err_cv_acc / 5;
78 end
79
80 %%
81 x_train = trainData;
82 x_train(:, d) = [];
83 y_train = trainData(:, d);
84 y_train = (y_train + 1) ./ 2;
85
86 x_test = testData;
87 x_test(:, d) = [];
88 y_test = testData(:, d);
89 y_test = (y_test + 1) ./ 2;
90 %%
91 for i = 1:6
92     d1 = d1_arr(i, 1);
93     w1_init = load(strcat(path_init, "w1_", num2str(d1), ".mat"));
94     w1_init = w1_init.w_1;
95
96     b1_init = load(strcat(path_init, "b1_", num2str(d1), ".mat"));
97     b1_init = b1_init.b1;
98
99     w2_init = load(strcat(path_init, "w2_", num2str(d1), ".mat"));
100     w2_init = w2_init.w_2;
101
102     b2_init = load(strcat(path_init, "b2_", num2str(d1), ".mat"));
103     b2_init = b2_init.b2;
104
105     w1 = w1_init;
106     b1 = b1_init;
107     w2 = w2_init;
108     b2 = b2_init;
```

Because writing ML code from scratch sucks (part 2)

```
107 w2 = w2_init;
108 b2 = b2_init;
109
110 for iter = 1: 5000
111     z = x_train * w1 + b1;
112     a = 1./(1 + exp(-z));
113     eta = 1./(1 + exp(-(a*w2+b2)));
114     eps = eta - y_train;
115     dgz = a .* (1 - a);
116
117     dw1 = (w2 .* ((eps .* dgz)' * x_train))'./m_train;
118     db1 = (w2 .* (dgz' * eps))'./m_train;
119     dw2 = (eps' * a)'./m_train;
120     db2 = sum(eps)/m_train;
121
122     w1 = w1 - step .* dw1;
123     b1 = b1 - step .* db1;
124     w2 = w2 - step .* dw2;
125     b2 = b2 - step .* db2;
126
127 end
128
129 pred_a_train = 1./(1 + exp(-(x_train * w1 + b1)));
130 pred_eta_train = 1./(1 + exp(-(pred_a_train * w2 + b2)));
131 pred_y_train = sign(pred_eta_train .* 2 - 1);
132 pred_y_train = (pred_y_train + 1) ./ 2;
133 err_train = classification_error(pred_y_train, y_train);
134 err_train_arr(1, 1) = err_train;
135
136 pred_a_test = 1./(1 + exp(-(x_test * w1 + b1)));
137 pred_eta_test = 1./(1 + exp(-(pred_a_test * w2 + b2)));
138 pred_y_test = sign(pred_eta_test .* 2 - 1);
139 pred_y_test = (pred_y_test + 1) ./ 2;
140 err_test = classification_error(pred_y_test, y_test);
141 err_test_arr(1, 1) = err_test;
142
143 end
144
145 %%
146 plot(d1_arr, err_train_arr, 'color', 'b'); hold on;
147 plot(d1_arr, err_test_arr, 'color', 'r'); hold on;
148 plot(d1_arr, err_cv_arr, 'color', 'g');
149
150 %%
151 disp(err_train_arr); hold on; %0.3960    0.1440    0.1120    0.0760    0.0640    0.0360
152 disp(err_test_arr); hold on; %0.3939    0.2066    0.1726    0.1563    0.1538    0.1544
153 disp(err_cv_arr); %0.3960    0.1840    0.1640    0.1000    0.1520    0.1440
154
155 %%
156 w1_init = load(strcat(path_init, "w1_", num2str(5), ".mat"));
```

```
160
161 b1_init = load(strcat(path_init, "b1_", num2str(5), ".mat"));
162 b1_init = b1_init.b1;
163
164 w2_init = load(strcat(path_init, "w2_", num2str(5), ".mat"));
165 w2_init = w2_init.w2;
166
167 b2_init = load(strcat(path_init, "b2_", num2str(5), ".mat"));
168 b2_init = b2_init.b2;
169
170 %%
171 x_train = trainData;
172 x_train(:, d) = [];
173 y_train = trainData(:, d);
174 y_train = (y_train + 1) ./ 2;
175
176 %%
177 x_test = testData;
178 x_test(:, d) = [];
179 y_test = testData(:, d);
180 y_test = (y_test + 1) ./ 2;
181
182 %%
183 step = 0.1;
184
185 w1 = w1_init;
186 b1 = b1_init;
187 w2 = w2_init;
188 b2 = b2_init;
189 for iter = 1: 5000
190     z = x_train * w1 + b1;
191     a = 1./(1 + exp(-z));
192     eta = 1./(1 + exp(-(a*w2+b2)));
193     eps = eta - y_train;
194     dgz = a .* (1 - a);
195
196     dw1 = (w2 .* ((eps .* dgz)' * x_train))'./m_train;
197     db1 = (w2 .* (dgz' * eps))'./m_train;
198     dw2 = (eps' * a)'./m_train;
199     db2 = sum(eps)/m_train;
200
201     w1 = w1 - step .* dw1;
202     b1 = b1 - step .* db1;
203     w2 = w2 - step .* dw2;
204     b2 = b2 - step .* db2;
205
206 end
207
208 pred_a_train = 1./(1 + exp(-(x_train * w1 + b1)));
209 pred_eta_train = 1./(1 + exp(-(pred_a_train * w2 + b2)));
210 pred_y_train = sign(pred_eta_train .* 2 - 1);
211 pred_y_train = (pred_y_train + 1) ./ 2;
```

Because writing ML code from scratch sucks (part 3)

```
209 pred_y_train = sign(pred_eta_train .* 2 - 1);
210 pred_y_train = (pred_y_train + 1) ./ 2;
211 err_train = classification_error(pred_y_train, y_train);
212 disp(err_train);
213
214 pred_a_test = 1./(1 + exp(-(x_test * w1 + b1)));
215 pred_eta_test = 1./(1 + exp(-(pred_a_test * w2 + b2)));
216 pred_y_test = sign(pred_eta_test .* 2 - 1);
217 pred_y_test = (pred_y_test + 1) ./ 2;
218 err_test = classification_error(pred_y_test, y_test);
219 disp(err_test);
220 %%
221 [w1, b1, w2, b2] = nn(x_train, y_train, w1_init, b1_init, w2_init, bw_init, step, iter);
222 %%
223 cv_train = load(strcat(path, num2str(1), "/cv-train.txt"));
224 cv_test = load(strcat(path, num2str(1), "/cv-test.txt"));
225
226 x_cv_train = cv_train;
227 x_cv_train(:, d) = [];
228 y_cv_train = cv_train(:, d);
229 x_cv_test = cv_test;
230 x_cv_test(:, d) = [];
231 y_cv_test = cv_test(:, d);
```


Contrast with PyTorch

```
37 # Fully connected neural network with one hidden layer
38 class NeuralNet(nn.Module):
39     def __init__(self, input_size, hidden_size, num_classes):
40         super(NeuralNet, self).__init__()
41         self.fc1 = nn.Linear(input_size, hidden_size)
42         self.relu = nn.ReLU()
43         self.fc2 = nn.Linear(hidden_size, num_classes)
44
45     def forward(self, x):
46         out = self.fc1(x)
47         out = self.relu(out)
48         out = self.fc2(out)
49         return out
```

```
51 model = NeuralNet(input_size, hidden_size, num_classes).to(device)
52
53 # Loss and optimizer
54 criterion = nn.CrossEntropyLoss()
55 optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
56
57 # Train the model
58 total_step = len(train_loader)
59 for epoch in range(num_epochs):
60     for i, (images, labels) in enumerate(train_loader):
61         # Move tensors to the configured device
62         images = images.reshape(-1, 28*28).to(device)
63         labels = labels.to(device)
64
65         # Forward pass
66         outputs = model(images)
67         loss = criterion(outputs, labels)
68
69         # Backward and optimize
70         optimizer.zero_grad()
71         loss.backward()
72         optimizer.step()
```


Deep learning frameworks

- Modern tools make it easy to implement neural networks
- Often used components
 - Linear, convolution, recurrent layers etc.
- Many frameworks available: Torch (2002), Theano (2011), Caffe (2014), TensorFlow (2015), PyTorch (2016)



Libraries for Deep Learning

- **Torch (Lua):**
 - <http://torch.ch/>
- **PyTorch (Python)**
 - <http://pytorch.org/>
- **TensorFlow (Python and C++):**
 - <https://www.tensorflow.org/>
- **Theano (Python)**
 - <http://deeplearning.net/software/theano/>
 - No longer maintained
- **Keras, PaddlePaddle, CNTK**

What is TensorFlow?

- Open source software library for numerical computation using data flow graphs
- Developed by Google Brain Team for machine learning and deep learning and made open-source
- TensorFlow provides an extensive suite of functions and classes that allow users to build various models from scratch

These slides are adapted from the following Stanford lectures:

https://web.stanford.edu/class/cs20si/2017/lectures/slides_01.pdf

<https://cs224d.stanford.edu/lectures/CS224d-Lecture7.pdf>

What's a tensor?

- **Formally, tensors are multilinear maps from vector spaces to the real numbers**
- **Think of them as n-dimensional array, with 0-d tensors being scalars, 1-d tensor vectors, 2-d tensor matrices, etc**

PyTorch

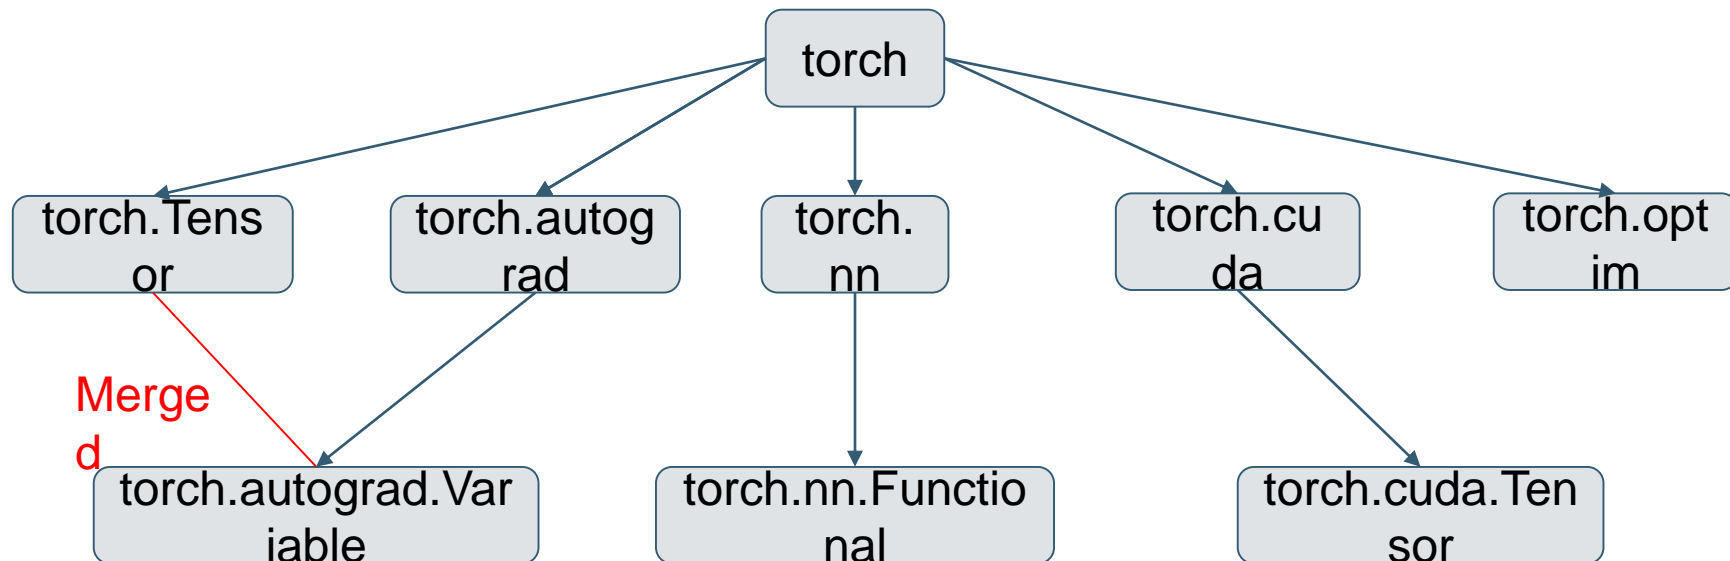
Pytorch

- Fast tensor computation (like numpy) with strong GPU support
- Deep learning research platform that provides maximum flexibility and speed
- Dynamic graphs and automatic differentiation

PyTorch

- 🔗 **Based on Torch, a scientific computing library for Lua**
- 🔗 **Developed by FAIR**
- 🔗 **Main features are the built-in computational graph and built-in GPU acceleration**

Structure of PyTorch library



Basic Concepts

torch.Tensor - similar to `numpy.array`

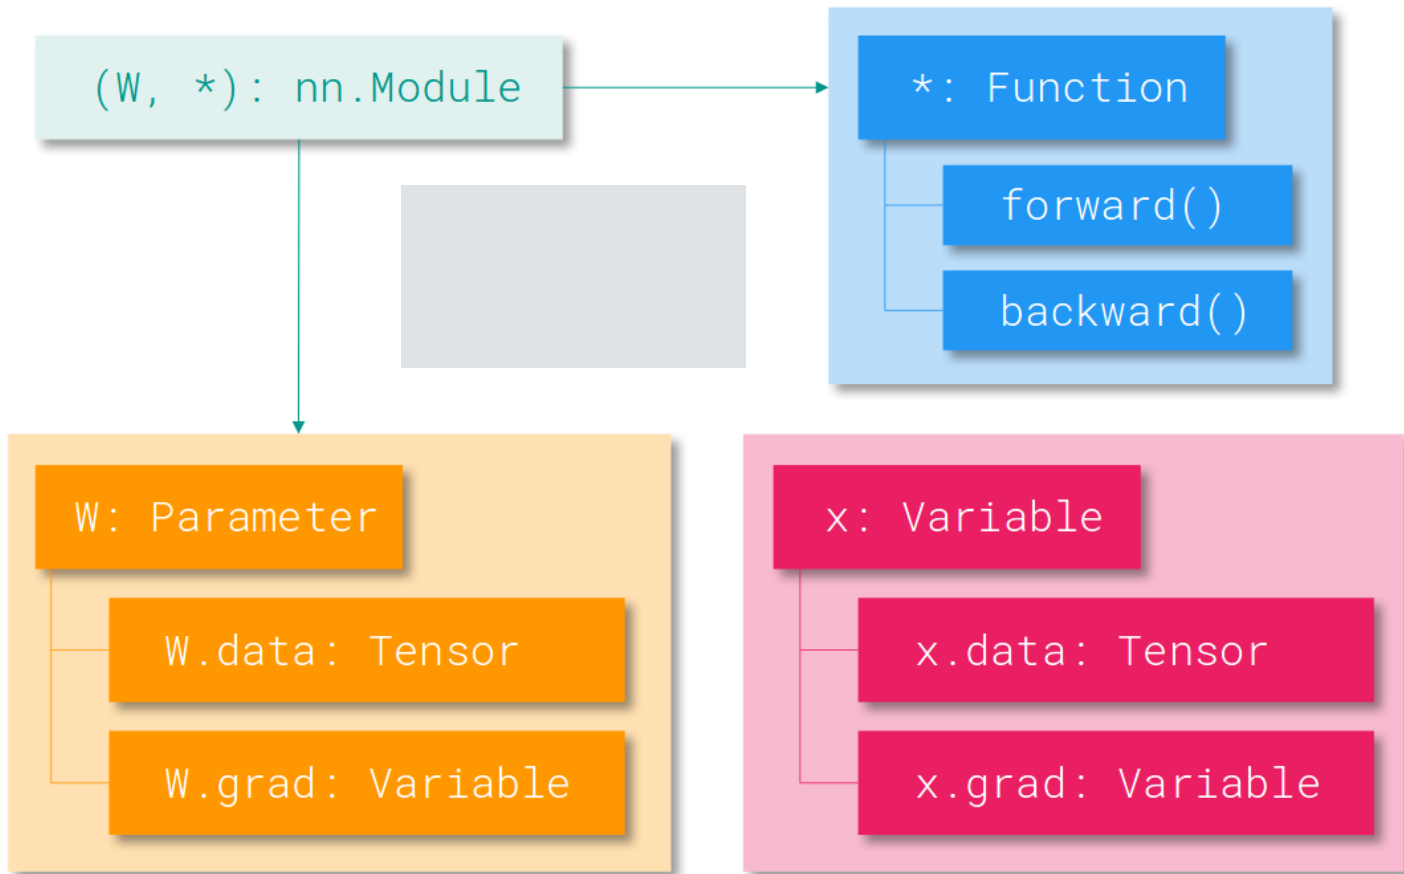
autograd.Variable - wraps a *Tensor* and enables auto differentiation

autograd.Function - operate on *Variables*, implements forward and backward

nn.Parameter - contain *Parameters* and define functions on input *Variables*

nn.Module - contain *Parameters* and define functions on input *Variables*

Basic Concepts



Basic Concepts

```
import torch  
x = torch.rand(10, 5)  
y = torch.rand(10, 5)  
z = x * y
```

Matrix Multiplication in Python

```
def matrixmult (A, B):  
    rows_A = len(A)  
    cols_A = len(A[0])  
    rows_B = len(B)  
    cols_B = len(B[0])  
  
    if cols_A != rows_B:  
        print "Cannot multiply the two matrices. Incorrect dimensions."  
        return  
  
    # Create the result matrix  
    # Dimensions would be rows_A x cols_B  
    C = [[0 for row in range(cols_B)] for col in range(rows_A)]  
    print C  
  
    for i in range(rows_A):  
        for j in range(cols_B):  
            for k in range(cols_A):  
                C[i][j] += A[i][k] * B[k][j]  
  
    return C
```

Matrix Multiplication in Numpy

```
>>> a=np.array([[1,2],[3,4]])  
>>> b=np.array([[11,12],[13,14]])  
>>> np.dot(a,b)  
array([[37, 40],  
       [85, 92]])
```


Defining a Neural Net in PyTorch

```
37 # Fully connected neural network with one hidden layer
38 class NeuralNet(nn.Module):
39     def __init__(self, input_size, hidden_size, num_classes):
40         super(NeuralNet, self).__init__()
41         self.fc1 = nn.Linear(input_size, hidden_size)
42         self.relu = nn.ReLU()
43         self.fc2 = nn.Linear(hidden_size, num_classes)
44
45     def forward(self, x):
46         out = self.fc1(x)
47         out = self.relu(out)
48         out = self.fc2(out)
49         return out
```

```
51 model = NeuralNet(input_size, hidden_size, num_classes).to(device)
52
53 # Loss and optimizer
54 criterion = nn.CrossEntropyLoss()
55 optimizer = torch.optim.Adam(model.parameters(), lr=learning_rate)
56
57 # Train the model
58 total_step = len(train_loader)
59 for epoch in range(num_epochs):
60     for i, (images, labels) in enumerate(train_loader):
61         # Move tensors to the configured device
62         images = images.reshape(-1, 28*28).to(device)
63         labels = labels.to(device)
64
65         # Forward pass
66         outputs = model(images)
67         loss = criterion(outputs, labels)
68
69         # Backward and optimize
70         optimizer.zero_grad()
71         loss.backward()
72         optimizer.step()
```


Network definition

```
class TwoLayerNet(torch.nn.Module):  
    def __init__(self, D_in, H, D_out):  
        super(TwoLayerNet, self).__init__()  
        self.linear1 = torch.nn.Linear(D_in, H)  
        self.linear2 = torch.nn.Linear(H, D_out)  
        self.relu = torch.nn.ReLU(inplace=True)  
  
    def forward(self, x):  
        h_relu = self.linear1(x)  
        h_relu = self.relu(h_relu)  
        y_pred = self.linear2(h_relu)  
    return y_pred
```





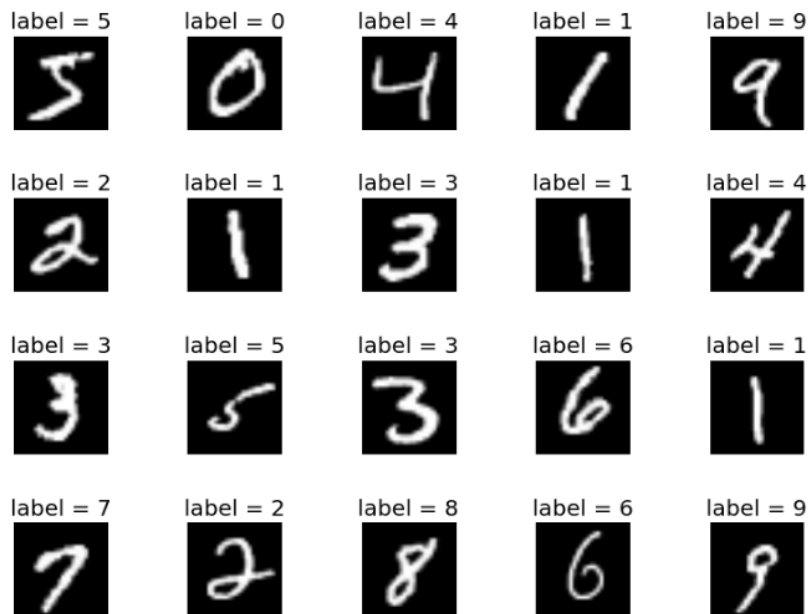
Optimize the Network

```
N, D_in, H, D_out = 64, 1000, 100, 10
model = TwoLayerNet(D_in, H, D_out)
loss_fn = torch.nn.MSELoss(size_average=False)
optimizer = torch.optim.SGD(model.parameters(), lr=1e-4)

for t in range(500):
    # Forward pass: Compute predicted y by passing x to the model
    y_pred = model(x)
    # Zero gradients, perform a backward pass, and update the weights.
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
```

Classifying Handwritten Digits





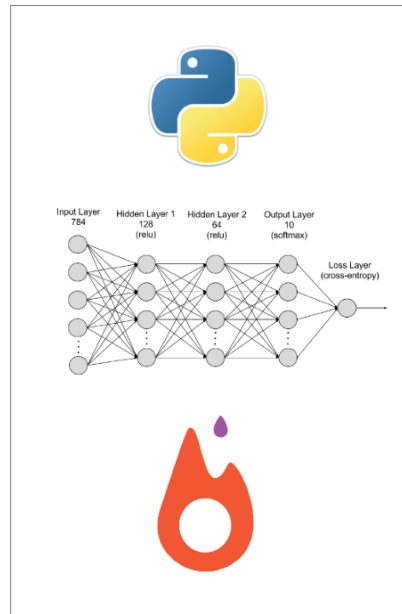
*is
the
new*

```
>>> print("Hello World!")  
Hello World!
```



0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
3 3 3 3 3 3 3 3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9 9 9 9 9 9 9 9

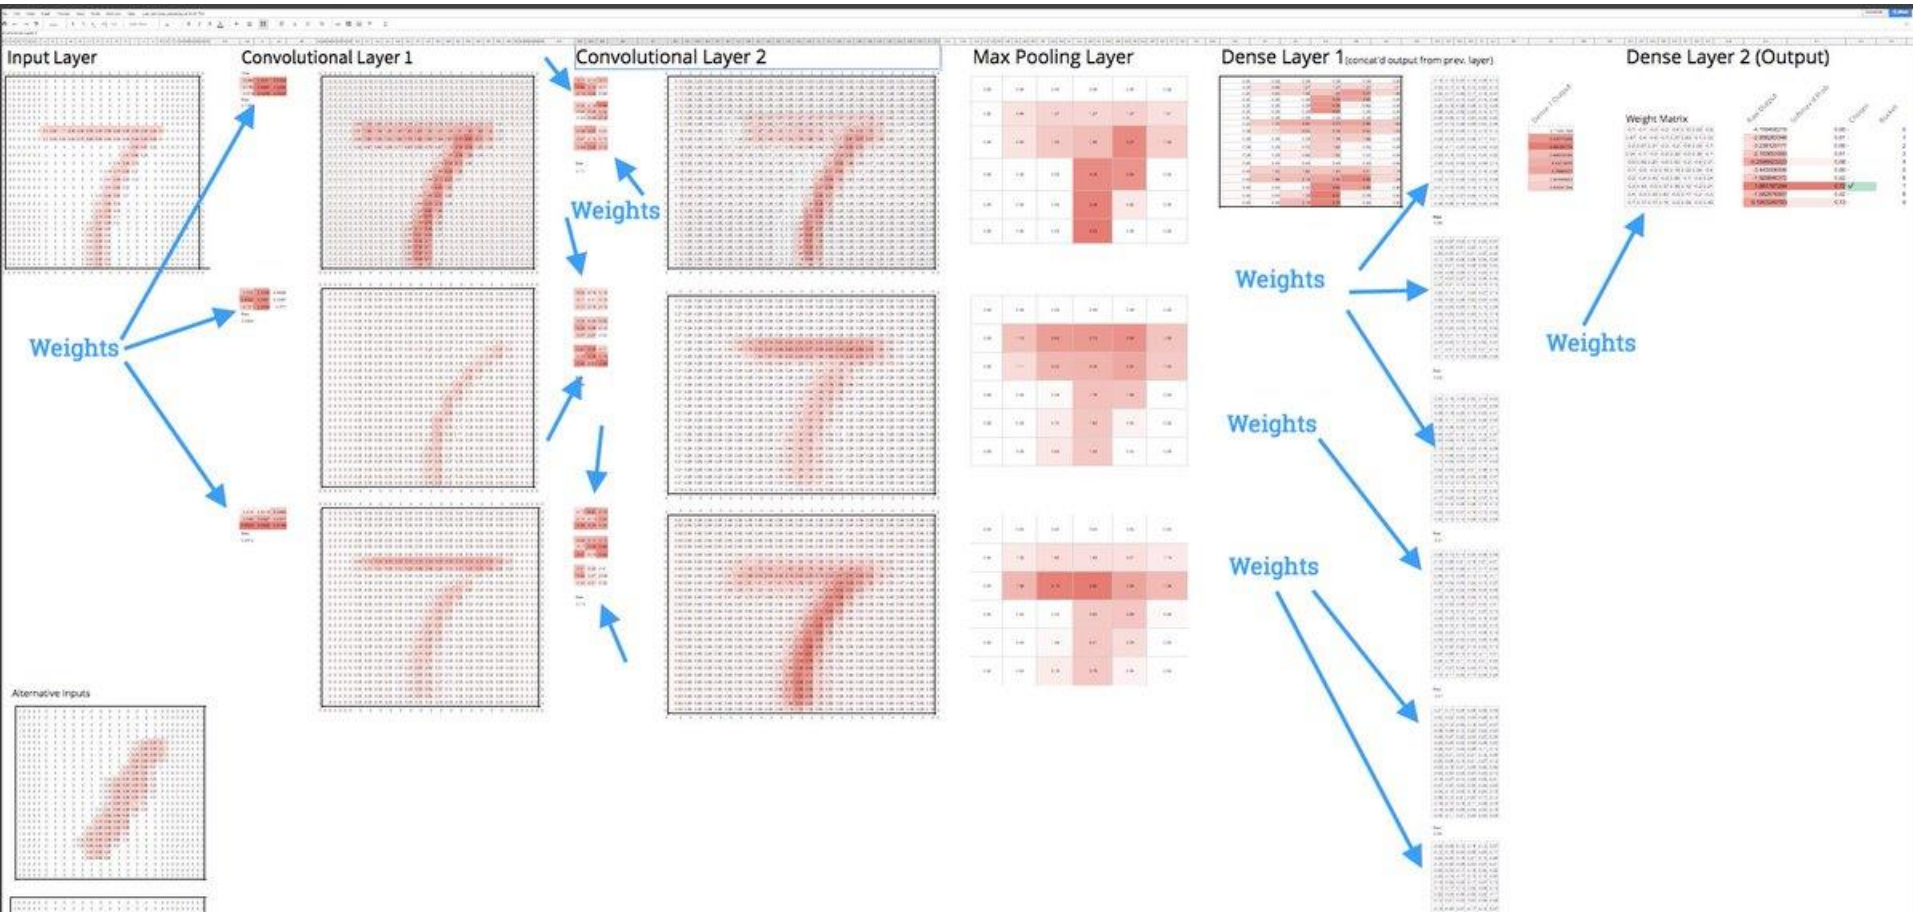


2



2





Define a CNN Network Class

```
class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        # 1 input image channel, 6 output channels, 5x5 square convolutional kernel
        self.conv1 = nn.Conv2d(1, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)
        # an affine operation: y = Wx + b
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        # Max pooling over a (2, 2) window
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        # If the size is a square you can only specify a single number
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

    def num_flat_features(self, x):
        size = x.size()[1:] # all dimensions except the batch dimension
        num_features = 1
        for s in size:
            num_features *= s
        return num_features
```

Questions?

- adnanmasood@gmail.com
- <https://twitter.com/adnanmasood>
- <https://www.linkedin.com/in/adnano>

Please use EventsXD to fill out a session evaluation.

Thank you!